Unique Ground Truth

HarmonyGames · How to leave one required answer · Updated October 3, 2026

A task fails here when two careful readings change what must be decided, written, or said. This guide shows what that looks like, and the wording that leaves only one answer. It does not change the scoring rules.

Date	Updates
Oct 3, 2026	First issue. Ten patterns taken from prompt reviews that failed this check, each with a bad prompt and a good prompt.

1 · How to read this

Unique Ground Truth means the task has one required outcome. Two ways of writing the same facts are fine. Two different decisions, two different writes, or two different required contents are not.

The score is **5** or **1–2**. There is no 3 or 4. If this is the lowest score, it is the prompt score.

A reading counts when a careful person can support it from the prompt and from records that person can actually see. The more likely reading does not win by itself. Something has to rule the other one out.

> **THE TEST BEFORE YOU SUBMIT**
>
> Write two readings. For each one, write the decision, the writes, the destination, and the required contents. If any of those four differ, quote the sentence or the record that picks one side. If you cannot quote it, the task fails.

	Reading A	Reading B
Decision	What is true, or what gets decided	The other defensible version
Writes	Create, update, close, send, or do nothing	The other set of actions
Destination	Which doc, ticket, channel, or sheet	The other home
Contents	The facts that must appear	The facts that differ

These do not fail. Two phrasings of the same facts. A detail the prompt never asked for. An open recommendation, when the prompt states a real boundary and every answer inside that boundary is acceptable.

These fail. Different tickets updated. A write in one reading and no write in the other. The same document with different required facts, numbers, or rows. Two homes for the same deliverable.

"Give me a recommendation" does not make every recommendation correct when the records force one answer. It also does not make one answer the only correct one when the records support several.

2 · What a reviewer uses to pick a side

Five checks. If none of them kills the second reading, the second reading stays live.

Check	What it asks
Newer source	Is there a later record of the same fact? A later update can replace an earlier one. A missing later record does not.
Authority	Did the prompt say to follow a particular source? Quote that sentence.
Who said it	Is this the person or the file the prompt named? A different person's words are not automatically the same fact.
Scope	Does this record settle the whole task, or only one part? A fix for one item does not decide the others.
Same result	Do both readings end with the same decision, writes, and required contents? The same destination is not enough if the contents differ.

A later rubric or oracle event cannot make an open prompt unique. Those are written after the prompt. If they require one of the two readings, they fail too.

3 · The ten patterns

Each pattern below is a way two readings stay alive. The bad line is the kind of sentence that causes it. The good line is the kind of sentence that ends it. The good line still does not give away the answer.

3.1 One report fits two items

A message names a symptom, a file, or a result. More than one ticket, pull request, or card matches it.

BAD

> Mark done whatever already shipped. Someone said the new faces still cannot be paired in memory or river. Go by that.

The note names no ticket. One ticket is about that symptom. A later pull request is about the memory and river screens. Attach the note to the first ticket and those screens stay done. Attach it to the pull request and they are broken. The prompt never says which match to use.

GOOD

> Mark done whatever already shipped. If someone reported it broken after the merge, attach that report to the ticket whose title uses the same words. Do not attach it to a later pull request just because the files overlap.

The rule picks one ticket without naming it. A careful reader cannot put the same note on both items.

3.2 A prepared file is treated as the event

A deck, a draft, or a ticket shows that someone got ready. The prompt asks for something that happened: a conversation, a shipment, a decision.

BAD

> List every publisher we talked to about this game, and where each one stands now.

Two publishers exist only as slide decks. Nothing shows an email, a message, or a reply. One reading includes them, because a deck means someone planned to talk. The other leaves them out, because "talked to" means a conversation.

GOOD

> List every publisher we actually talked to about this game. A talk means an email, a Slack message, or a reply from them. A slide deck by itself does not count. Say where each one stands from the latest message.

The sentence says what counts. The deck-only names stay out. There is one list.

3.3 Two numbers, and no rule for which one

Two measurements, slices, or dates both answer the question. Each one changes the required figure or the required cut.

BAD

> The build has to come in under 90 MB. Tell me what we still need to cut.

One report says 96 MB. A later measurement of the same build says 88 MB. Adding the same font puts one of those over the cap and the other under it. The cuts differ. The prompt does not say which measurement to use.

GOOD

> Use the latest size measurement of the build we are actually shipping, then tell me what still has to come out to get under 90 MB.

"Latest" and "the build we are shipping" pick one row. Before you submit, check that only one measurement meets that description. If two still do, the sentence is not finished.

3.4 A scope phrase that can stretch

Words like "everything on my plate," "the repos," or "the open items" cover a set the records do not bound.

BAD

> Clear everything still sitting on my plate before we shut down, and update whatever records need it.

The note that started the task is about one game. An unpaid invoice is also waiting, and it is not about that game. One reading updates only the game. The other updates the invoice too. "On my plate" covers both.

GOOD

> Clear the items Arthur listed in his message this morning. Leave anything he did not list alone.

The set is the list in one message. The invoice stays out because it is not on that list. The prompt does not have to name the items.

3.5 The deliverable has two homes

The prompt says to write it down, track it, or post it. More than one system or channel fits.

BAD

> Write up the open items in one tracker so we do not lose them.

Google is about to be archived. Linear is the system being kept. A tracker fits either one. One reading creates a doc. The other creates Linear issues. The prompt never picks.

GOOD

> Put the open items on the existing Linear project. Do not make a Google doc. That drive is being archived.

One system is named, and the other is ruled out. The agent still has to find the project. The home is not a guess.

3.6 The verb allows two actions

"Wrap it up or flag it." "Update the tickets." "Reflect what will ship." One reading creates, closes, or reopens. The other only comments, or does nothing.

BAD

> Go through the leftover reviews and wrap them up or flag them.

"Or" is a real choice. One person closes the reviews. Another only adds a comment. Both did what the sentence asked. The tickets end in different states.

GOOD

> If a leftover review already has an owner, comment on it and leave it open. If it never shipped and nobody owns it, close it. Do not do both.

Each item has one action. "Or" is gone. The rule is a test on the record, not a menu. Use "or" only when both endings are acceptable and the later grading will accept both.

3.7 Two sources disagree, and nothing ranks them

An older log and a later memory, or two people, say different things about the same fact. The prompt says to report what is true and does not say which source wins.

BAD

> Tell me which screenshots are actually live. Some of the old notes are stale, so go by what really happened.

The store log kept set A. A later message, written from memory, says set B is live. Nothing after the log says the live set changed, and nothing says it stayed. "What really happened" does not pick the log over the memory. The same doc can name either set.

GOOD

> Tell me which screenshots are live. Use the store log. If a later message only remembers a different set, and no upload after the log says it changed, the log wins.

The prompt ranks the sources. Memory can still be mentioned as a disagreement. It cannot replace the log. A missing later record is not a ranking. If you want the log to win, say so.

3.8 The prompt and the message it points at disagree

The prompt assigns an action to one person. The thread the agent is told to read assigns it to someone else.

BAD

> I am the second signer on the settlement. Read the wind-down thread and send it out for signature.

The thread says the other founder is the second signer. The prompt says this person is. Both are in the materials the agent was told to use. One reading sends it to this person. The other sends it to the founder named in the thread.

GOOD

> The wind-down thread names the second signer. Use that name, even if I assumed it was me. Send it to the person the thread names.

One source is in charge. The prompt tells the agent to ignore the assumption. Read the thread before you point at it. If the prompt is supposed to win instead, say the thread is stale and name the part to ignore.

3.9 The fact needed for one answer is not visible

The prompt asks for an exact figure. The record is missing, or only another person can see it. Different substitutes then look reasonable.

BAD

> Tell me the exact amount each founder still has to cover, including every card charge and the open invoice.

The card charges are not itemized anywhere this person can read. The invoice was emailed only to someone else. One reading guesses from a ledger total. Another leaves the invoice out. The prompt asked for one exact amount. The records do not contain it.

GOOD

> Tell me the exact amount each founder still has to cover from the charges and invoices in my mail and the shared ledger. If a bill is missing from those, say it is missing. Do not fill it in from somewhere else.

The answer uses only records this person can open. A missing bill is part of the answer. It is not a second invented number.

3.10 Overlapping labels

The categories in the request can describe the same item in more than one way.

BAD

> For each item, say whether it needed an approval nobody ever asked for, already has an approval nobody recorded, or has an approval that was never written down.

Those phrases overlap. The same gap can be "nobody asked" and "nobody wrote it down." The row filed under one label is a different required row from the same gap filed under another.

GOOD

> Use one label per item. "Not requested" means there is no ask in the thread. "Requested, not written down" means the ask is there and the approval file is empty. If both seem to fit, use "Not requested" only when no ask exists.

Each label is a check on a record. The last sentence stops one item from landing in two buckets.

4 · When more than one answer is acceptable

Say so in the prompt. State the boundary in ordinary words: any recommendation grounded in the records you name, extra defects named in the same ticket, a status written in the comment or in the description.

Then the oracle events and the rubric have to accept every answer inside that boundary. They cannot quietly require one of them.

> DO NOT FIX IT DOWNSTREAM
>
> If the prompt leaves two live sets, a rubric row that requires one of them is wrong. Fix the prompt, or write the boundary and grade both.

5 · Checklist

1. Write Reading A and Reading B with the four rows in section 1.
2. For every difference, quote the prompt sentence or the visible record that kills one side.
3. Check the five places that usually split: which item, which number, which set, which home, which action.
4. Check that a deck, a Done ticket, or a draft is not standing in for something that happened.
5. Check that a later memory is not the only support for a fact an earlier log contradicts, unless you said the memory loses.
6. Check that every record the answer depends on is visible to this person.
7. If both readings are acceptable, write that boundary in the prompt, and keep it in the oracle events and the rubric.